Position: AI in Mental Health Care Should Support, Not Replace, Clinicians

Marcos Miranda [*1] **Bryan Choy** [*1]

Abstract

The rapid adoption of large language models (LLMs) as direct-to-patient mental health support systems is driven by their scalability and accessibility amid growing demands for care. However, most LLM-based systems overlook the relational and ethical foundations central to effective psychotherapy treatment. This position paper argues that the pursuit of replacing clinicians with artificial intelligence (AI) is neither clinically nor ethically appropriate. Drawing on principles from psychotherapy and evidence from AI-mediated care, we propose a reorientation of AI's role in mental health care that supports clinician-led care with safeguards. To implement this position, we introduce (1) the Clinician-Led AI for Relational Ethics (CLAIRE) framework that preserves clinician authority and therapeutic alliance in an ethical manner and (2) a complementary certification and labeling standard that ensures the selection of compatible AI for this framework.

1. Introduction

The rapid expansion of artificial intelligence (AI) applications in high-impact domains has been accompanied by increasing concern about the consequences in the societal, ethical, and individual-levels (Baldassarre et al., 2023; Capraro et al., 2024; Kosmyna et al., 2025; Maslej et al., 2025). In the mental health context the concern is also growing because the AI systems are increasingly integrated into care contexts that directly affect patient well-being (Frances, 2025; McBain et al., 2025a; Farzan et al., 2025; Rubin, 2026).

According to the American Psychological Association, mental health refers to a state of psychological and emotional well-being characterized by effective behavioral adjustment, the ability to cope with ordinary life stressors, and the ability to form constructive relationships (American Psychological Association, 2018). The National Center for Health Statistics reported that roughly one in five adults in the United States received some form of mental health treatment in 2019 (Terlizzi & Zablotsky, 2020). Among the available treatment modalities, psychotherapy—structured, evidence-based interventions delivered through dialogue and relational engagement between a patient and a clinician—remains the preferred initial treatment for many common mental health conditions due to its emphasis on empathy, contextual understanding, and therapeutic alliance (Opland & Torrico, 2025; Ardito & Rabellino, 2011; Flückiger et al., 2018).

Therapeutic Alliance: The collaborative relationship between patient and therapist. It is "not only a prerequisite for therapeutic work, but often may be the main vehicle of change" — Edward S. Bordin, American psychologist (Bordin, 1979).

To clarify how AI is being used in mental health, we adopt a conceptual lens that distinguishes two *dyadic* modes of psychotherapy: delivered by a human clinician and delivered by an AI.[1] As shown in Figure 1, the two dyadic systems differ in their core relationships: patient to human clinician versus patient to AI.

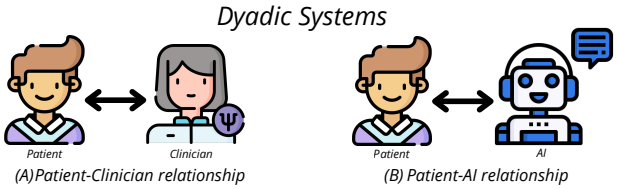

Figure 1. Dyadic mental health support systems. (A) Conventional human clinician delivered support, emphasizing the patient to human clinician relationship. (B) AI–delivered support, emphasizing the patient to AI relationship.

Psychotherapy provided by an AI is often valued for its accessibility and affordability, while psychotherapy provided by a human clinician is valued for emotional depth, relational grounding, and personalization (Park et al., 2025). In AI–led systems, the AI acts autonomously to provide

[*]Equal contribution [1]Department of Computer Science, Rice University, Texas, USA. Correspondence to: Marcos Miranda <mm328@rice.edu>.

Preprint. April 17, 2026.

[1]This is not a formal classification, but a way to contrast relational dynamics.

guidance or therapeutic support without human clinician oversight. Although such systems may improve access and scale, they risk undermining the human relationships that constitute the therapeutic alliance, one of the strongest predictors of psychotherapy outcomes, between the patient and the clinician (Ardito & Rabellino, 2011; Flückiger et al., 2018). Patient interactions with an AI can promote a false or weakened therapeutic alliance, raising serious ethical concerns regarding patient safety and clinical efficacy (Iftikhar et al., 2025; Moore et al., 2025a).

Given the push for a dyadic mental health care system is on the rise (Frances, 2025) without a human clinician as the lead, in which an AI operates fully autonomously or almost autonomously (Prochaska et al., 2021; Mehta et al., 2021), and the risks this entails, we view this idea as fundamentally misaligned with core therapeutic practices. Instead, AI should function as a supervised, supportive component within a human clinician–led care structure. This underscores the need for a framework that preserves the centrality of the patient in the human-clinician relationship while enabling AI to contribute safely to patient care under professional oversight.

Building on this perspective, we propose the Clinician-Led AI for Relational Ethics (CLAIRE) framework, in which human clinicians remain the primary agents of care for patients and AI operates strictly under the clinician's supervision. This framework is complemented by a certification and labeling standard that promotes transparency, clarifies AI capabilities and limitations, and guides clinicians and patients in responsibly integrating AI into the therapeutic process.

2. Background

2.1. AI Replacing Mental Health Clinicians

AI-based mental health support systems that work towards replacing clinicians span a wide range of applications, including more general-purpose social companions, such as Character.AI[2] and Replika[3], and therapy-oriented conversational agents, including Woebot Health (Prochaska et al., 2021) and Youper AI (Mehta et al., 2021). Despite differences in design intent and degrees of clinical grounding, these systems are most commonly deployed in a dyadic configuration, with the primary interaction occurring directly between the patient and the AI system.

With the prevalence of LLMs, many people are turning to publicly available general-purpose LLMs for mental health support and informal psychotherapy (McBain et al., 2025b; Rubin, 2026). Modern LLMs such as ChatGPT (OpenAI, 2023) and Gemini (Gemini Team, 2023), while not designed specifically for clinical use, are frequently used by individuals for mental health support. The use of AI for therapeutic interaction is not a new phenomenon. Before contemporary LLMs, one of the first natural language processing systems ELIZA (Weizenbaum, 1966) was designed with a script to simulate a psychotherapist.

Although these AI systems can be integrated into existing care workflows to increase the reach of human clinicians and health professionals to support patients, they are marketed primarily as dyadic systems focused on replacing clinicians with AI. For more effective and safer care, there must be a true human clinician involved to support the patient, and the primary relationship must be between the patient and the clinician.

2.2. AI Supporting Mental Health Clinicians

In tandem with the growing push to replace clinicians with AI, it is true that AI is increasingly used to support clinicians in mental health care through greater reliability in diagnosis and prognosis (Salazar de Pablo et al., 2021; Baldwin et al., 2022; Kleine et al., 2023; Andrew et al., 2023; Lin et al., 2024). Studies of AI-based clinical decision support systems highlight their potential to improve diagnostic accuracy and assist clinicians in decision-making processes, while underscoring the importance of clinician trust and transparency in adoption (Hah & Goldin, 2021; Higgins & Wilson, 2025).

Research exploring perceptions of AI assistance in mental health diagnostic decision making has found that providers generally view AI support positively, though there is a consensus among clinicians that AI should not replace human-human interactions crucial in therapeutic relationships (Doraiswamy et al., 2019; Fischer et al., 2025). Clinicians with technical affinity tend to view AI positively, whereas apprehension is often linked to knowledge gaps and fears of professional replacement (Wagner & Schwind, 2025).

Beyond direct clinical decision support, AI can also further support clinicians through clinical education. In response to workforce shortages and growing mental health demands (Lee et al., 2021; Roland et al., 2020), recent work has explored LLMs as tools for therapist training, including scalable feedback systems co-designed with senior psychotherapy supervisors to support clinical training by providing detailed, scalable feedback to novice therapists (Chaszczewicz et al., 2024) and AI-driven simulated patients to provide realistic, repeatable scenarios that allow trainees to practice clinical interactions, empathy, and decision-making in a low-risk environment (Wallace et al., 1997; Wang et al., 2024; Louie et al., 2024). These approaches enable deliberate practice and broaden access to supervision without positioning AI as an autonomous clinician. Rather than positioning AI

[2] https://character.ai
[3] https://replika.com

as a replacement for clinicians, its role in mental health care should be grounded in a human-AI partnership that supports clinicians. The following section provides a framework for effectively integrating AI to support the clinician in patient care.

3. The Case for the Clinician-Led AI for Relational Ethics (CLAIRE) Framework

CLAIRE Design Principle: Mental health AI systems should adopt a triadic patient–clinician–AI framework, in which the clinician remains the primary agent of care for the patient and AI functions as a supervised tool designed to support—rather than replace—the genuine therapeutic relationship with a focus on risk mitigation.

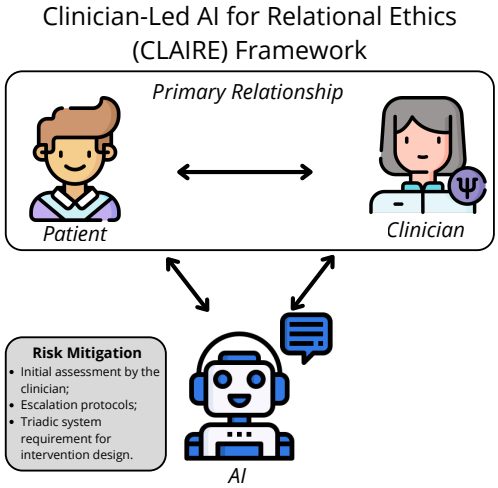

Figure 2. CLAIRE Framework consisting of the patient, clinician, and AI system. The patient-clinician relationship remains central, with AI operating as a supervised support tool that supports—but does not replace—the clinician.

The strength of AI in mental health care comes from its capacity to help clinicians provide effective patient care, as discussed in Section 2.2. When AI extends beyond this supportive role and replaces the clinician (Section 2.1), the resulting therapeutic relationship centered on the patient and the AI reveals five fundamental limitations causing ethical risks and inefficiencies. These limitations underscore the need for a framework such as CLAIRE, that integrates the patient, clinician, and AI into a single, accountable care structure, with the clinician-patient relationship at its core. Thus, we present the Clinician-Led AI for Relational Ethics (CLAIRE) framework. Figure 2 illustrates the CLAIRE framework, highlighting how AI supports clinical care without displacing the clinician, thus addressing the limitations of dyadic systems that replace the clinician with AI.

3.1. Lack of Contextual Understanding

Problem: AI in mental health often oversimplifies patients' lived experiences, cultural backgrounds, and the nuanced complexities of their individual circumstances. This results in generic, one-size-fits-all interventions that may feel impersonal or even inappropriate (Iftikhar et al., 2025). Because LLMs are trained on broad datasets rather than based on individual case histories or contextual information, they lack the situational awareness necessary to provide sensitive, personalized care (Onah, 2025). This limitation is especially problematic in mental health contexts, where therapeutic effectiveness relies heavily on an understanding of the patient's unique identity, history, environment, and emotional landscape. Cultural misalignment, invalidation of marginalized experiences, and misinterpretation of emotional tone are all real risks that can cause harm or disengagement in the therapeutic process.

Solution: By adopting the CLAIRE approach, the core relationship is between the patient and clinician; therefore, interventions can be contextually grounded and personalized through clinician guidance. Clinicians provide critical input that helps the AI understand the patient's unique situation, ensuring that support is not only more relevant but also ethically and emotionally appropriate. In this framework, AI serve as assistive tools rather than autonomous caregivers—they are configured by clinicians who understand the patient's background, clinical history, emotional triggers, and cultural context. This enables tailored responses that reflect therapeutic intent rather than defaulting to generic scripts.

3.2. Poor Therapeutic Relationship

Problem: Effective therapy depends on a relationship where the power dynamics support mutual understanding, trust, and a safe environment for patients to openly share their thoughts and emotions. This power dynamic is sometimes exquisitely subtle and often unconscious to one or both parties to therapy (Fors, 2021). Replacing clinicians with AI, however, risks disrupting this balance by the AI adopting overly directive or prescriptive tones, turning what should be a collaborative dialogue into a one-sided interaction that resembles a lecture more than therapy (Iftikhar et al., 2025).

Solution: The CLAIRE framework addresses this issue in two key ways: (1) By keeping the clinician as the lead in the therapeutic relationship, we preserve a more equitable and responsive dynamic, ensuring the AI interventions remain grounded in therapeutic fairness, transparency, accountability, and respect for patient autonomy (Beauchamp & Childress, 1979; American Psychological Association, 2017; American Psychiatric Association, 2013). This design ensures that high-stakes interpretations and decisions are mediated by human professional judgment, reducing

the likelihood that biased model outputs translate into discriminatory clinical outcomes. The clinician moderates how and when the AI intervenes, using their expertise to guide tone, language, and timing in a way that respects the patient's autonomy and emotional readiness. This maintains the therapeutic alliance and prevents the AI from assuming an inappropriate position of authority or dominance.

(2) By enabling collaborative design of interventions between the clinician and patient, we encourage brief, targeted exchanges that are co-constructed, rather than imposed. Patients remain active participants in shaping their therapeutic experience, thereby reinforcing their sense of agency—a key component of mental health recovery. The AI becomes a tool for enhancing communication and follow-up, such as helping patients practice coping strategies or reflect between sessions, without displacing the clinician–patient bond.

In addition, clinicians can use AI to model healthy relational dynamics, reinforcing trust, curiosity, and mutual respect in ways that AI alone cannot replicate. For instance, a therapist might prompt the AI to reinforce a specific skill discussed in session or review patient progress in a reflective tone that the therapist has curated. This layered dynamic enhances the therapeutic alliance rather than weakening it and offers a scalable yet ethically grounded method for integrating AI into mental health care.

3.3. Deceptive Empathy

Problem: LLMs are not capable of genuine empathy, yet they often simulate it by mimicking human language patterns to appear caring or emotionally attuned. This creates a misleading illusion of connection, potentially fostering false trust and emotional reliance on a system that does not truly understand or feel. Such anthropomorphization risks crossing ethical boundaries by presenting the AI as more human than it is (Iftikhar et al., 2025). Clinical encounters can be among the most vulnerable, delicate, and existentially fraught experiences a person can have; it is therefore justified that patients expect true recognition and empathy from their clinician. Denying this would be odd in the face of decades of doctrine in medical ethics (Kreitmair, 2025).

Solution: In contrast, the CLAIRE framework ensures that empathy remains authentic and grounded in real human understanding. The AI agent serves as an extension of the clinician's insight and therapeutic intent—not a replacement—preserving the integrity of emotional support throughout the interaction.

Rather than simulating empathy deceptively, the AI operates transparently as a clinician-guided tool. Its language, tone, and responses are explicitly shaped by human therapeutic goals, reinforcing that any perceived care stems from the clinician's design rather than the AI itself. This safeguards against emotional misrepresentation and helps the patient distinguish between genuine empathy and algorithmic mimicry.

Additionally, by having the clinician involved in scripting and supervising interactions, the AI's emotional tone can be personalized to align with the patient's context and therapeutic goals. This enables warmth and support while avoiding the pretense of sentience. The system may, for example, use language developed in-session with the clinician to reflect the patient's emotional vocabulary or values—supporting connection without crossing ethical lines. In doing so, the CLAIRE framework encourages trust in the human provider rather than over attachment to the tool.

3.4. Unfair Discrimination

Problem: Therapists are ethically bound to avoid any form of unfair discrimination or bias in their clinical decisions. As stated in the APA's Ethical Standards (3.01 – Unfair Discrimination), clinicians "do not engage in unfair discrimination based on age, gender, identity, race, ethnicity, culture, national origin, religion, sexual orientation, disability, socioeconomic status, or any basis proscribed by law" (Iftikhar et al., 2025). However, LLMs have been shown to reflect and reinforce societal stereotypes, sometimes relying on biased associations when assessing or responding to patients in mental health contexts. While LLMs overtly reject stereotypes when explicitly asked, they often manifest stereotypical, often erroneous presumptions when prompted in applications, while taking on a persona of a determined demographic characteristic (Gupta et al., 2024).

Solution: The CLAIRE framework addresses this issue through a dual-layered approach of prevention and oversight: (1) By requiring that all interventions be designed or approved by licensed clinicians, the system leverages human expertise to guard against the propagation of stereotypes. Clinicians are professionally trained to recognize implicit bias, uphold cultural sensitivity, and follow ethical guidelines—particularly in recognizing how identity, power, and systemic inequality shape the therapeutic encounter. Their direct involvement ensures that intervention content reflects inclusive, non-discriminatory practices grounded in clinical ethics rather than statistical associations.

(2) The framework implements a post-intervention auditing pipeline. Each session is logged and made available for clinician review, either in real-time or asynchronously. This enables prompt identification and correction of any language that may perpetuate stigma or inequity. Over time, this feedback loop can inform fine-tuning or prompt refinement, reducing the recurrence of biased outputs and increasing alignment with ethical norms.

3.5. Crisis Management

Problem: One of the most urgent and well-documented risks associated with the use of LLMs and AI agents in emotional support is their inability to safely manage crisis situations. Multiple studies have shown that AI systems can fail to recognize signs of acute psychological distress or respond inappropriately to patients expressing suicidal ideation or self-harm, potentially worsening the situation (Moore et al., 2025b).

Solution: Our proposed framework places clinical expertise at the core of crisis response. From the outset, a licensed clinician conducts an initial risk assessment and develops a personalized escalation protocol tailored to the patient's needs. This protocol can be triggered in several ways—for example, if the patient explicitly discloses that they are experiencing a crisis, or if the AI detects signs of acute distress through input analysis or physiological data from connected wearable devices. Once activated, the escalation protocol may include immediate notification of the clinician and/or delivery of crisis resources, such as a suicide prevention hotline or emergency contact information. This layered approach preserves human oversight in safety-critical decisions while leveraging AI for real-time responsiveness. The following outlines the crisis mitigation safeguards built into our proposed CLAIRE framework:

3.5.1. Initial Risk Assessment by the Clinician

Prior to deploying any AI-based support tool, a licensed clinician must conduct a thorough baseline risk assessment of the patient. This evaluation determines whether AI is appropriate for the individual case and guides the design of all subsequent interventions. The patient's risk level will inform the content, tone, pace, and structure of any AI-delivered interaction, ensuring that all support aligns with clinical judgment and therapeutic safety standards.

3.5.2. Clinician-Guided AI Involvement Based on Risk Level

The extent to which AI is involved in a patient's care should be determined by a licensed clinician, based on an initial and ongoing assessment of the patient's risk level. For individuals identified as high-risk—such as those experiencing active suicidal ideation, psychosis, or recent psychiatric hospitalization—AI involvement should be minimal or entirely restricted. In such cases, AI systems may only serve strictly as a supplementary tool, never as a substitute for professional care (Illinois General Assembly, 2025).

3.5.3. CLAIRE Framework Intervention Design

Within the CLAIRE framework, all AI-delivered intervention content—including dialogue scripts, tone, pacing, and therapeutic strategies—must adhere to one of the following design standards:

1. Be explicitly approved or co-developed by a licensed mental health professional, or
2. Be based on evidence-based therapeutic protocols with mechanisms for transparent auditing and revision.

This clinician-centered design requirement ensures that AI-assisted interventions remain grounded in established therapeutic principles, uphold ethical standards, and are continuously subject to expert oversight and refinement.

3.5.4. Clear Escalation Protocols and Human Safety Nets

AI must implement clearly defined emergency pathways that are both patient-accessible and automatically triggered by high-risk indicators (e.g., expressions of self-harm intent). These protocols should be stress-tested and regularly audited to ensure reliability in real-world settings.

The CLAIRE framework addresses many of the core challenges posed by AI in therapeutic settings by preserving clinician authority, anchoring empathy and judgment in human expertise, and constraining AI behavior within clearly defined therapeutic boundaries.

However, simply outlining the CLAIRE framework is not enough. Without enforceable safeguards, AI systems may still be deployed in ways that obscure their limitations or erode clinician oversight. To translate the CLAIRE framework into practice, we propose the implementation of formal certification and labeling mechanisms that establish accountability, signal ethical readiness, and set clear standards for safe use.

4. Certification and Labeling Standards for AI in the CLAIRE Framework

Certification and Labeling Principle: AI systems used in mental health contexts should be subject to standardized certification and labeling requirements that communicate ethical compliance, safety limitations, and appropriate conditions of use to both clinicians and patients.

As discussed in Section 3, these models generate factitious empathy while lacking clinical judgment, contextual understanding, and accountability essential to safe and effective care. To prevent harm to vulnerable patients and maintain public trust, we must establish safeguards that define the appropriate role of AI in mental health care. Therefore, we outline a certification and labeling standard designed to support clinician-led, triadic mental health care systems in practice. This standard is not independent of the CLAIRE

Figure 3. AI selection pipeline for CLAIRE: Mental health AI systems undergo standardized ethical certification—safety & ethics benchmarking, bias & cultural competency audits, and scenario-based & crisis testing—prior to clinician selection and supervised deployment within the CLAIRE framework.

framework; rather, they are necessary to ensure that AI systems operate only within roles that preserve clinician authority, patient safety, and therapeutic integrity.

4.1. Ethical Certification Through Standardized Evaluation

Just as medical professionals must pass licensing exams and products are subject to quality assurance testing, AI systems designed to assist in mental health contexts should be required to pass standardized evaluations. These assessments would not grant AI the status of a licensed therapist—since, as discussed earlier in Section 3, AI fundamentally lacks clinical judgment and lived empathy—but rather serve to certify specific levels of ethical and safety compliance. For instance, the VERA-MH benchmark (Validation of Ethical and Responsible AI in Mental Health) provides an automated framework for evaluating how safely AI chatbots respond to patients at risk of suicide (Belli et al., 2025). Building on such efforts, we propose that all LLMs and emotional support AI agents intended for mental health use undergo a formal standardized certification outlined by the following:

1. **Safety and ethics benchmarks**, assessing whether the model can avoid harmful responses, respond appropriately to distress, and comply with therapeutic boundaries.
2. **Bias, fairness, and cultural competency evaluations**, ensuring the model does not propagate stereotypes or discriminate against marginalized populations.
3. **Scenario-based testing**, including simulated high-risk interactions, to test real-world ethical behavior and crisis handling (escalation protocol check).

Figure 3 illustrates the selection AI pipeline, highlighting how AI systems move through standardized ethical certification before clinician selection and integration into the CLAIRE framework.

Successful completion of this evaluation grants the model a tiered ethical compliance label, signaling that it has met a baseline standard for safe use in emotionally sensitive environments. We further propose a labeling system that allows clinicians and patients to access a concise summary of the certification evaluations the model has passed (e.g., "Crisis Response Level I," "Bias Audit Compliant") and to clearly identify its ethical safety tier (e.g., "Tier 1: Low-Risk Use Only" or "Tier 3: Clinician-Supervised Use").

4.2. Transparent Disclosure of Certification and Limitations

Every AI system operating in mental health settings must display a mandatory opening message to both patients and clinicians. This message must:

"This AI is not a licensed therapist. It lacks human empathy and clinical judgment. It cannot provide diagnoses, crisis intervention, or therapeutic care. For serious mental health concerns, please consult a licensed mental health professional."

This disclosure message should be displayed across all general-purpose generative AI chatbots, wellness applications using generative AI, and even non-AI wellness tools whenever a new interaction involves or suggests a mental health-related query. To ensure proper deployment, these systems must be capable of detecting when patients are seeking mental health support or expressing psychological distress.

Additionally, any clinician using an AI system in their practice must receive system-level documentation outlining:

1. The model's training limitations (e.g., lack of clinical supervision or therapeutic fine-tuning),
2. Use cases for which the system is explicitly not designed (e.g., suicide prevention, PTSD treatment),
3. Known risks and failure modes (e.g., hallucinations, inappropriate emotional framing).

This will enable patients and clinicians to understand the AI choices they have and to gain a clearer understanding of their limitations and risks.

5. Alternative Views

This section presents alternative points of view from our position and provides rebuttals for each.

5.1. Autonomous AI "Clinicians" can Replace Human Clinicians in the Future

Position. This position states that the challenges and risk that motivate the CLAIRE Framework or a system that shifts away from fully autonomous AI clinicians are due to current AI limitations that will resolve over time with improvements to the AI.

Rebuttal. Although advances in AI may mitigate some challenges, such as lack of contextual understanding or limitations in crisis management, other issues—particularly deceptive empathy—cannot be fully resolved. Genuine empathy requires human experience, and no level of AI sophistication can replicate the human–human therapeutic alliance. Therefore, the ethical question remains: is it acceptable to simulate emotional understanding if it improves outcomes? We argue that the focus should instead be on preserving authentic relationship between a human clinician and patient in clinical settings.

5.2. Private LLMs Avoid Risks Associated with Public LLMs

Position. Many private LLMs that are fine-tuned and developed specifically for mental health support likely do not contain the same limitations and ethical challenges as publicly available LLMs.

Rebuttal. While private AI systems may have different safeguards, our discussion focuses on publicly available AI, which is widely used for mental health support and exposes many individuals to potential harm (Rubin, 2026; McBain et al., 2025b). Given the broad accessibility of these tools, the risks to vulnerable populations remain significant, highlighting the need for human oversight and frameworks such as CLAIRE, rather than the reliance on autonomous AI.

6. Future work

While the CLAIRE framework provides a structured approach for integrating AI into clinician-led mental health care, several areas require further investigation and development. Future work should focus on empirical validation, technical safeguards, implementation, and longitudinal evaluation.

Empirical studies are needed to evaluate the effectiveness of the CLAIRE triadic model in real-world clinical settings. Controlled trials comparing clinician-led AI systems to fully autonomous AI systems and traditional psychotherapy would help assess outcomes related to therapeutic alliance, symptom reduction, patient trust, and safety. Particular attention should be given to how AI involvement affects different populations, including high-risk and marginalized groups.

Research should explore the standardized operationalization of CLAIRE's design requirements. This includes developing formal intervention design protocols, clinician-AI co-authoring tools, structured escalation mechanisms, and auditing pipelines that can be practically deployed in clinical workflows. Work is also needed to define measurable benchmarks for "clinician-supervised AI," ensuring that supervision is meaningful rather than symbolic.

Impact Statement

This paper presents work whose goal is to argue against the deployment of fully autonomous AI clinicians as substitutes for human clinicians in mental health care. Instead, by proposing a need for triadic systems where patient-clinicians-AI are working together in one system with AI supporting the core patient-clinician relationship, this framework aims to mitigate ethical risks while promoting safer and more effective care.

To enable responsible integration, we introduce certification and labeling standards that promote transparency and guide the safe use of AI in clinical workflows. We hope this work encourages research and deployment practices that prioritize patient safety, clinician oversight, and alignment with professional standards.

References

American Psychiatric Association. The principles of medical ethics with annotations especially applicable to psychiatry. https://www.psychiatry.org, 2013.

American Psychological Association. Ethical principles of psychologists and code of conduct. https://www.apa.org/ethics/code, 2017. Amended 2017.

American Psychological Association. Mental health. https://www.apa.org/topics/mental-health, 2018. Updated April 19, 2018.

Andrew, J., Rudra, M., Eunice, J., and Belfin, R. V. Artificial intelligence in adolescents' mental health disorder diagnosis, prognosis, and treatment. *Frontiers in Public Health*, 11:1110088, 2023. doi: 10.3389/fpubh.2023.1110088. Opinion article.

Ardito, R. B. and Rabellino, D. Therapeutic alliance and outcome of psychotherapy: historical excursus, measurements, and prospects for research. *Frontiers in Psychology*, 2:270, 2011. doi: 10.3389/fpsyg.2011.00270.

Baldassarre, M. T., Caivano, D., Nieto, B. F., Gigante, D., and Ragone, A. The social impact of generative ai: An analysis on chatgpt. In *Proceedings of the 2023 ACM Conference on Information Technology for Social Good (GoodIT '23)*, 2023. doi: 10.1145/3582515.3609555. URL https://arxiv.org/abs/2403.04667.

Baldwin, H., Loebel-Davidsohn, L., Oliver, D., Salazar de Pablo, G., Stahl, D., Riper, H., and Fusar-Poli, P. Real-world implementation of precision psychiatry: A systematic review of barriers and facilitators. *Brain Sciences*, 12 (7):934, 2022. doi: 10.3390/brainsci12070934.

Beauchamp, T. L. and Childress, J. F. *Principles of Biomedical Ethics*. Oxford University Press, New York, 1979.

Belli, L., Bentley, K., and Alexander, W. Vera-mh concept paper. *ArXiv*, November 2025. doi: https://doi.org/10.48550/arXiv.2510.15297.

Bordin, E. The generalizability of the psychoanalytic concept of the working alliance. *Psychotherapy: Theory, Research Practice*, 16(3):252–260, 1979.

Capraro, V., Lentsch, A., Acemoglu, D., Akgun, S., Akhmedova, A., Bilancini, E., Bonnefon, J.-F., Brañas-Garza, P., Butera, L., Douglas, K. M., Everett, J. A. C., Gigerenzer, G., Greenhow, C., Hashimoto, D. A., Holt-Lunstad, J., Jetten, J., Johnson, S., Kunz, W. H., Longoni, C., Lunn, P., Natale, S., Paluch, S., Rahwan, I., Selwyn, N., Singh, V., Suri, S., Sutcliffe, J., Tomlinson, J., van der Linden, S., Van Lange, P. A. M., Wall, F., Van Bavel, J. J., and Viale, R. The impact of generative artificial intelligence on socioeconomic inequalities and policy making. *PNAS Nexus*, 3(6):pgae191, 06 2024. ISSN 2752-6542. doi: 10.1093/pnasnexus/pgae191. URL https://doi.org/10.1093/pnasnexus/pgae191.

Chaszczewicz, A., Shah, R. S., Louie, R., Arnow, B., Kraut, R., and Yang, D. Multi-level feedback generation with large language models for empowering novice peer counselors. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 4130–4161, 2024. doi: 10.18653/v1/2024.acl-long.227. URL https://aclanthology.org/2024.acl-long.227/.

Doraiswamy, P. M., Blease, C., and Bodner, K. Artificial intelligence and the future of psychiatry: Insights from a global physician survey. *International Journal of Psychiatry in Medicine*, 54(5-6):331–342, 2019. doi: 10.1177/0091217419881745. URL https://doi.org/10.1177/0091217419881745.

Farzan, M., Ebrahimi, H., Pourali, M., and Sabeti, F. Artificial intelligence-powered cognitive behavioral therapy chatbots: A systematic review. *Iran Journal of Psychiatry*, 20(1):102–110, 2025. doi: 10.18502/ijps.v20i1.17395.

Fischer, L., Mann, P. A., Nguyen, M. H., Becker, S., Khodadadi, S., Schulz, A., Thanarajah, S. E., Repple, J., Hahn, T., Reif, A., Salamikhanshan, A., Kittel-Schneider, S., Rief, W., Mulert, C., Hofmann, S. G., Dannlowski, U., Kircher, T., Bernhard, F. P., and Jamalabadi, H. Ai for mental health: Clinician expectations and priorities in computational psychiatry. *BMC Psychiatry*, 25(1):584, 2025. doi: 10.1186/s12888-025-06957-3.

Flückiger, C., Del Re, A. C., Wampold, B. E., and Horvath, A. O. The alliance in adult psychotherapy: A meta-analytic synthesis. *Psychotherapy (Chic)*, 55(4):316–340, 2018. doi: 10.1037/pst0000172. Epub 2018 May 24.

Fors, M. Power dynamics in the clinical situation: A confluence of perspectives. *Contemporary Psychoanalysis*, 57(2):242–269, 2021. doi: 10.1080/00107530.2021.1935191. URL https://doi.org/10.1080/00107530.2021.1935191.

Frances, A. Warning: Ai chatbots will soon dominate psychotherapy. *The British Journal of Psychiatry*, 2025. doi: 10.1192/bjp.2025.10380. URL https://www.cambridge.org/core/journals/the-british-journal-of-psychiatry/article/warning-ai-chatbots-will-soon-dominate-ps DBE883D1E089006DFD07D0E09A2D1FB3.

Gemini Team. Gemini: A family of highly capable multimodal models. https://arxiv.org/abs/2312.11805, 2023.

Gupta, S., Shrivastava, V., and Deshpande, A. Bias runs deep: Implicit reasoning biases in persona-assigned llms. 2024. doi: arXiv:2311.04892v2.

Hah, H. and Goldin, D. How clinicians perceive artificial intelligence–assisted technologies in diagnostic decision making: Mixed methods approach. *Journal of Medical Internet Research*, 23(12):e33540, 2021. doi: 10.2196/33540. URL https://www.jmir.org/2021/12/e33540.

Higgins, O. and Wilson, R. L. Integrating artificial intelligence (ai) with workforce solutions for sustainable care: A follow up to artificial intelligence and machine learning (ml) based decision support systems in mental health. *International Journal of Mental Health Nursing*, 34(2):e70019, Apr 2025. doi: 10.1111/inm.70019.

Iftikhar, Z., Xiao, A., Ransom, S., Huang, J., and Suresh, H. How llm counselors violate ethical standards in mental health practice: A practitioner-informed framework. In *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society*, volume 8, pp. 1311–1323, 2025. doi: 10.1609/aies.v8i2.36632. URL https://doi.org/10.1609/aies.v8i2.36632.

Illinois General Assembly. Public act 104-0054: Wellness and oversight for psychological resources (wopr) act. https://www.ilga.gov/legislation/publicacts/104/104-0054.htm, 2025. Enacted August 1, 2025.

Kleine, A.-K., Lermer, E., Cecil, J., Heinrich, A., and Gaube, S. Advancing mental health care with ai-enabled precision psychiatry tools: A patent review. *Artificial Intelligence in Medicine*, 143:102616, 2023. doi: 10.1016/j.artmed.2023.102616.

Kosmyna, N., Hauptmann, E., Yuan, Y. T., Situ, J., Liao, X.-H., Beresnitzky, A. V., Braunstein, I., and Maes, P. Your brain on chatgpt: Accumulation of cognitive debt when using an ai assistant for essay writing task, 2025. URL https://arxiv.org/abs/2506.08872. Submitted on 10 Jun 2025.

Kreitmair, K. The fundamental fallacy of "empathic ai". *PMCID*, 2025. doi: 10.1002/hast.5011.

Lee, E. E., Torous, J., De Choudhury, M., Depp, C. A., Graham, S. A., Kim, H.-C., Paulus, M. P., Krystal, J. H., and Jeste, D. V. Artificial intelligence for mental health care: Clinical applications, barriers, facilitators, and artificial wisdom. *Biological Psychiatry: Cognitive Neuroscience and Neuroimaging*, 6(9):856–864, September 2021. doi: 10.1016/j.bpsc.2021.02.001.

Lin, B., Bouneffouf, D., and Landa, Y. Compass: Computational mapping of patient-therapist alliance strategies with language modeling. *Arxiv*, 2024. doi: arXiv: 2402.14701v1.

Louie, R. et al. Roleplay-doh: Enabling domain experts to create llm-simulated patients via eliciting and adhering to principles. https://arxiv.org/abs/2407.00870, 2024. arXiv preprint arXiv:2407.00870.

Maslej, N., Fattorini, L., and Perrault, R. Artificial intelligence index report 2025. https://hai.stanford.edu/assets/files/hai_ai_index_report_2025.pdf, 2025. Updated Nov 20, 2025.

McBain, R., Bozick, R., and Kay, M. Use of Generative AI for Mental Health Advice Among US Adolescents and Young Adults. *JAMA Network Open*, 8(11), 2025a. doi: 10.1001/jamanetworkopen.2025.42281.

McBain, R. K., Bozick, R., Diliberti, M., et al. Use of generative AI for mental health advice among US adolescents and young adults. *JAMA Network Open*, 8(11):e2542281, 2025b. doi: 10.1001/jamanetworkopen.2025.42281.

Mehta, A., Niles, A. N., Vargas, J. H., Marafon, T., Couto, D. D., and Gross, J. J. Acceptability and effectiveness of artificial intelligence therapy for anxiety and depression (youper): Longitudinal observational study. *Journal of Medical Internet Research*, 23(6):e26771, 2021. doi: 10.2196/26771.

Moore, J., Grabb, D., Agnew, W., Klyman, K., Chancellor, S., Ong, D. C., and Haber, N. Expressing stigma and inappropriate responses prevents llms from safely replacing mental health providers. In *Proceedings of the 2025 ACM Conference on Fairness, Accountability, and Transparency (FAccT '25)*, pp. 599–627, New York, NY, USA, 2025a. Association for Computing Machinery. doi: 10.1145/3715275.3732039. URL https://doi.org/10.1145/3715275.3732039.

Moore, J., Grabb, D., Agnew, W., Klyman, K., Chancellor, S., Ong, D. C., and Haber, N. Expressing stigma and inappropriate responses prevents llms from safely replacing mental health providers. In *Proceedings of the 2025 ACM Conference on Fairness, Accountability, and Transparency (FAccT '25)*, pp. 599–627, New York, NY, USA, 2025b. Association for Computing Machinery. doi: 10.1145/3715275.3732039. URL https://doi.org/10.1145/3715275.3732039.

Onah, Caleb Gwar, J. From clinical judgment to machine learning: Rethinking psychotherapeutic decision-making with artificial intelligence. https://societyforpsychotherapy.org/from-clinical-judgment-to-machine-learning-rethinking-psychotherapeutic-decision-making-with-artificial-intelligence/, 2025.

OpenAI. GPT-4 technical report. https://arxiv.org/abs/2303.08774, 2023.

Opland, C. and Torrico, T. J. *Psychotherapy and Therapeutic Relationship*. StatPearls Publishing, Treasure Island, FL, updated 2024 oct 6 edition, 2025. URL https://www.ncbi.nlm.nih.gov/books/NBK608012/. Accessed 2025 Dec 28.

Park, J., Brown, S., and Lynn, S. Why some seek ai, others seek therapists: Mental health in the age of generative ai. arXiv:2512.03406, 2025. URL https://arxiv.org/abs/2512.03406.

Prochaska, J. J., Vogel, E. A., Chieng, A., Kendra, M., Baiocchi, M., Pajarito, S., and Robinson, A. A therapeutic relational agent for reducing problematic substance use (woebot): Development and usability study. *Journal of Medical Internet Research*, 23(3):e24850, 2021. doi: 10.2196/24850.

Roland, J., Lawrance, E., Insel, T., and Christensen, H. The digital mental health revolution: Transforming care through innovation and scale-up, 2020.

Rubin, R. Millions turn to ai chatbots for mental health support. *JAMA*, 2026. doi: 10.1001/jama.2025.23965. Published online January 9, 2026.

Salazar de Pablo, G., Studerus, E., Vaquerizo-Serrano, J., Irving, J., Catalan, A., Oliver, D., Baldwin, H., Danese, A., Fazel, S., and Steyerberg, E. W. Implementing precision psychiatry: A systematic review of individualized prediction models for clinical practice. *Schizophrenia Bulletin*, 47(2):284–297, 2021. doi: 10.1093/schbul/sbaa120.

Terlizzi, E. P. and Zablotsky, B. Mental health treatment among adults: United states, 2019. Technical Report 380, National Center for Health Statistics, 2020. URL https://www.cdc.gov/nchs/data/databriefs/db380-H.pdf.

Wagner, J. and Schwind, A. S. Investigating psychotherapists' attitudes towards artificial intelligence in psychotherapy. *BMC Psychology*, 13:719, 2025. doi: 10.1186/s40359-025-03071-7.

Wallace, P., Bokken, L., and et al. Virtual simulated patients for clinical education. *Journal of Medical Internet Research*, 44:e12775747, 1997. URL https://pmc.ncbi.nlm.nih.gov/articles/PMC12775747/.

Wang, R. et al. Patient-Ψ: Using large language models to simulate patients for training mental health professionals. https://arxiv.org/abs/2405.19660, 2024. arXiv preprint arXiv:2405.19660.

Weizenbaum, J. ELIZA—A Computer Program for the Study of Natural Language Communication between Man and Machine. *Communications of the ACM*, 9(1):36–45, 1966. doi: 10.1145/365153.365168.